



FlexNN: Efficient and Adaptive DNN Inference on Memory-Constrained Edge Devices

Xiangyu Li¹, Yuanchun Li^{1,2}, Yuanzhe Li¹, Ting Cao³, Yunxin Liu^{1,2}

¹Institute for AI Industry Research (AIR), Tsinghua University

²Shanghai Artificial Intelligence Laboratory ³Microsoft Research

Edge AI is transforming our lives

The capability of AI has been demonstrated on various tasks.

Edge AI is taking an important role, but its capability is limited.

- Local DNN inference: network-free, cost-efficient, privacy-preserving.
- Challenge: resource constraints on edge devices.



Autonomous Vehicle



Smartphone



Edge Camera

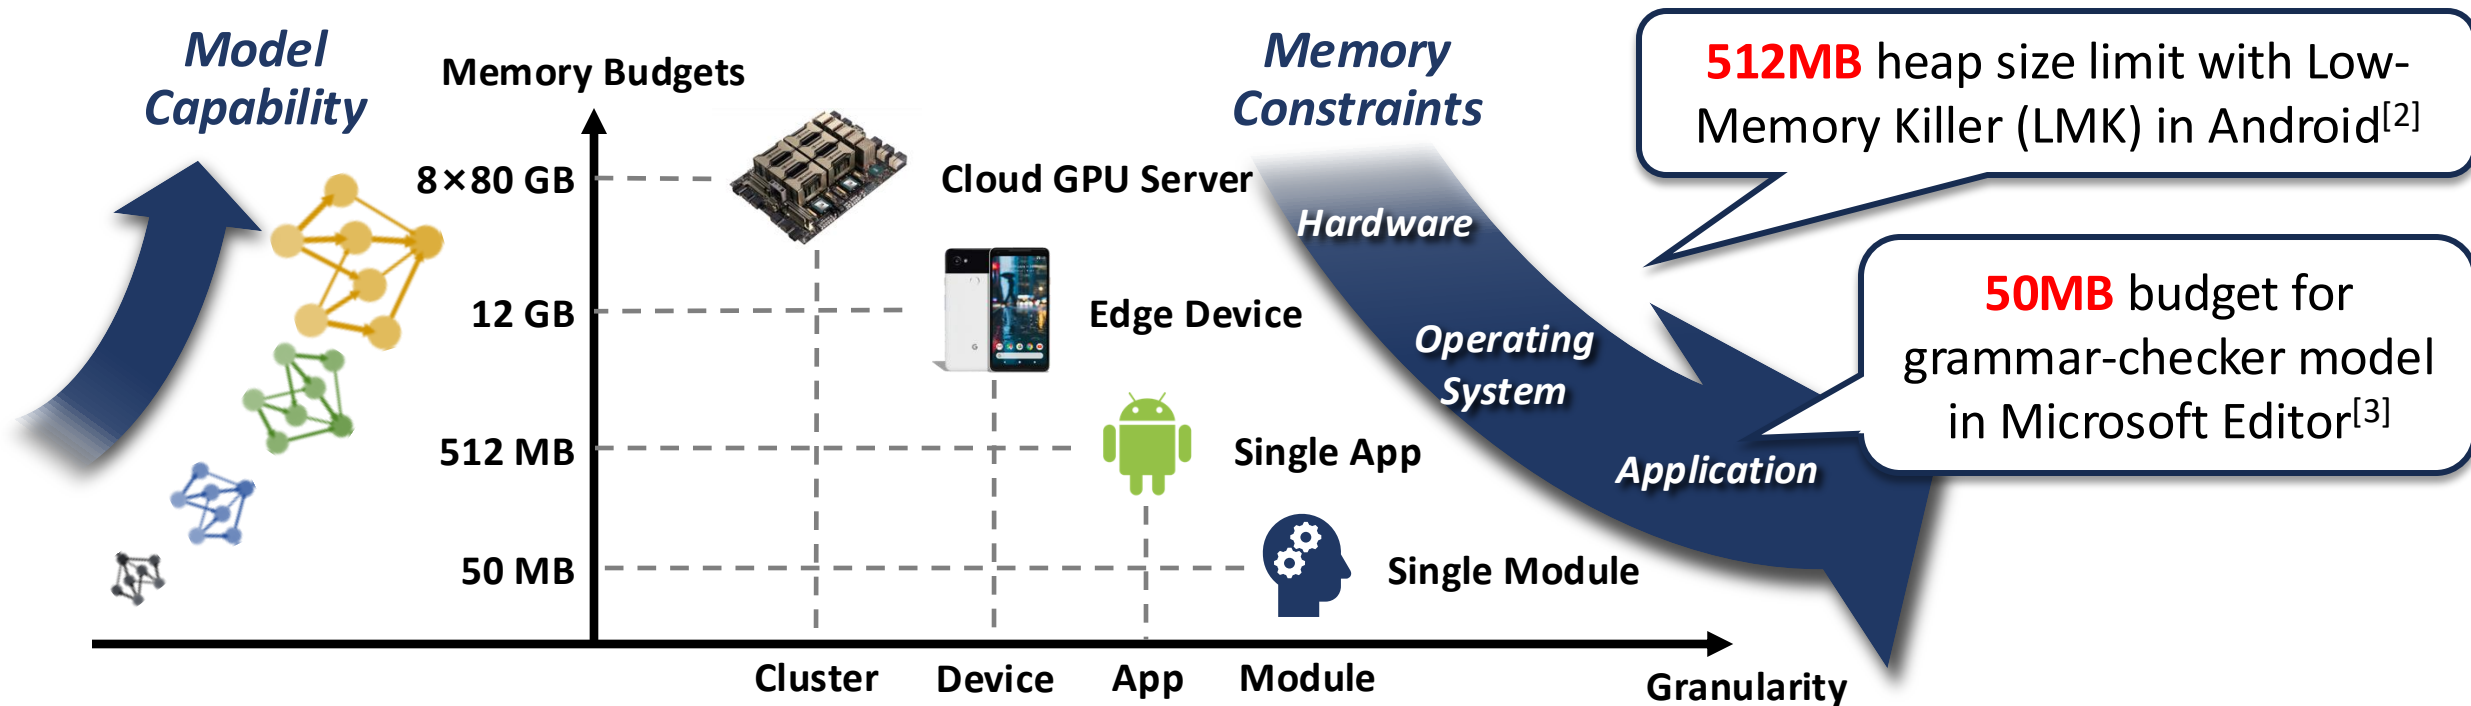
Pictures from Internet.

Memory bottleneck of on-device DNN inference

The improvement of memory has greatly lagged behind computation.

- From iPhone 4 to iPhone 14^[1]: GPU (1000×) > CPU (100×) >> RAM (8×)

There are multiple levels of memory constraints besides hardware.



[1] GadgetVersus. <https://gadgetversus.com/smartphone/apple-iphone-4-vs-appleiphone-14/>, accessed: 2023-08-01.

[2] Android Developers. <https://developer.android.com/topic/performance/memory>, accessed: 2023-08-01.

[3] Ge Tao, et al. <https://www.microsoft.com/en-us/research/blog/achievingzero-cogs-with-microsoft-editor-neural-grammar-checker/>, accessed: 2023-08-11.

Existing approaches for memory reduction

Model Customization

- Model compression (*e.g.*, quantization, pruning, ...)
- Efficient structure design (*e.g.*, MobileNets^[1])
- Neural Architecture Search



System Design


- Mobile inference frameworks (*e.g.*, NCNN^[4], MNN^[5], ...)
- System Optimizations
 - **layer streaming** (*i.e.*, layer-wise swapping)
 - **Layer partitioning** in TEE-based inference^[3]
 - **Memory planning** in on-device training^[2]

← deployment efforts ↑
model accuracy ↓

← limited optimization

← cannot directly apply

challenges



[1] Mark Sandler, et al. "Mobilenetv2: Inverted residuals and linear bottlenecks." CVPR 2018.

[2] Qipeng Wang, et al. "Melon: Breaking the memory wall for resource-efficient on-device machine learning." MobiSys 2022.

[3] Taegyeong Lee, et al. "Occlumency: Privacy-preserving remote deep-learning inference using SGX." MobiCom 2019.


[4] Hui Ni, and The ncnn contributors. Ncnn. 2017, <https://github.com/Tencent/ncnn>.

[5] Xiaotang Jiang, et al. "MNN: A universal and efficient inference engine." MLSys 2020.

Challenge 1: unbalanced memory footprints

Only a few layers have large memory footprints!

- In *layer streaming*, the “largest” layers determine the peak memory.
- They become memory bottlenecks of the whole model.



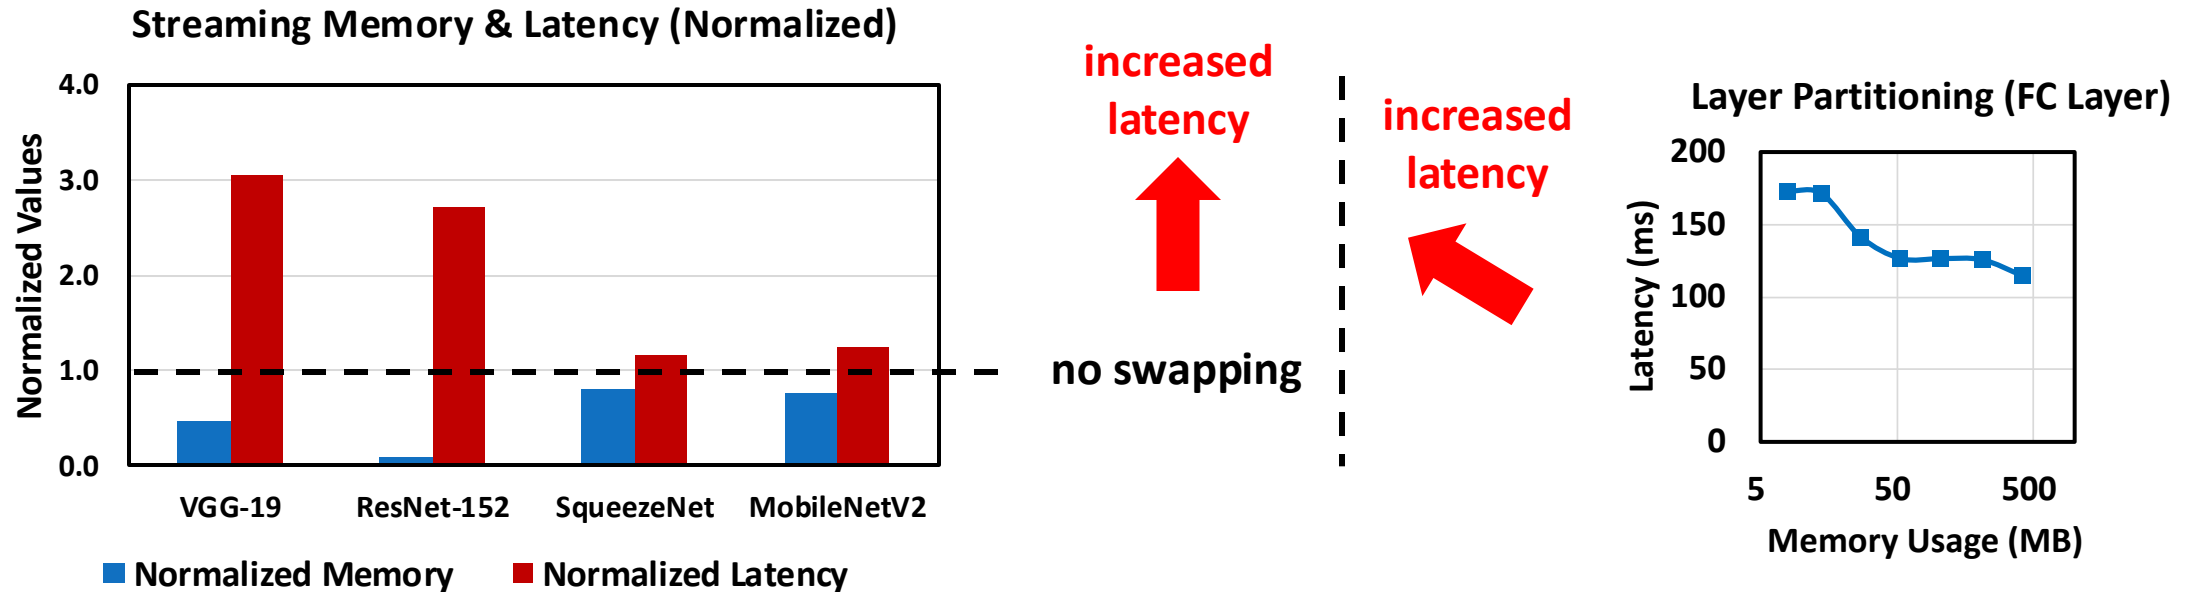
Layer-wise memory footprint of DNNs. Only a few layers have large memory footprints.

Layer IDs are consistent with the converted NCNN model format.

Challenge 2: memory management overhead

Memory reduction requires extra operations.

- **Layer streaming:** additional I/O and processing for repetitive weight loading.
- **Layer partitioning:** splitting/merging overhead and more memory fragments.



Memory and Latency Impacts of Naïve Layer Streaming.

The values are normalized to those without swapping.

Memory and Latency Trade-off of Naïve Layer Partitioning.

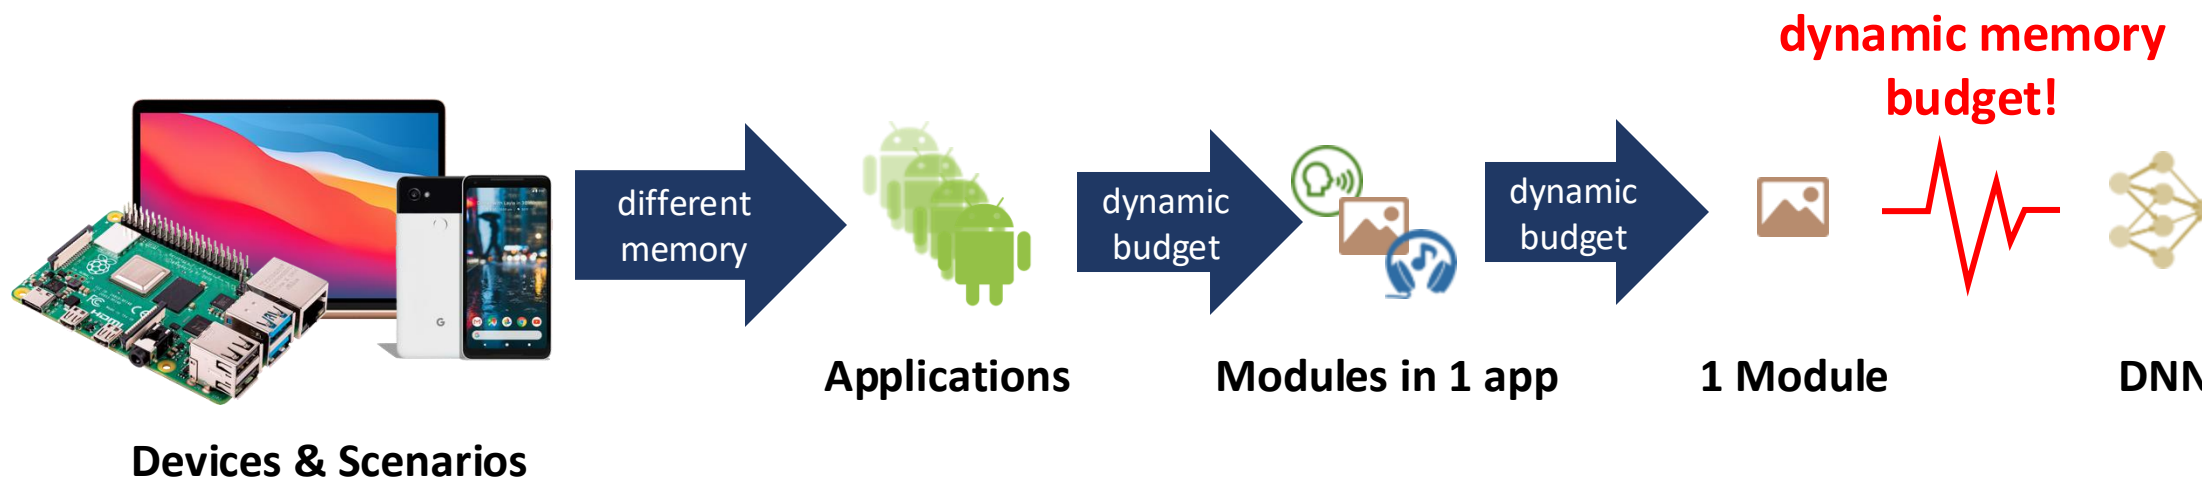
Challenge 3: memory budget dynamicity

In real-world applications, the memory budget can frequently change.

- Diverse devices and scenarios.
- Influences from other apps and modules.

It is important to fully utilize the memory while not exceeding the budget!

- Existing memory management schemes cannot address this issue.



Our approach: fine-grained joint planning

Challenges

Unbalanced memory distribution

Memory management overhead

Memory budget dynamicity

Goals

memory

latency

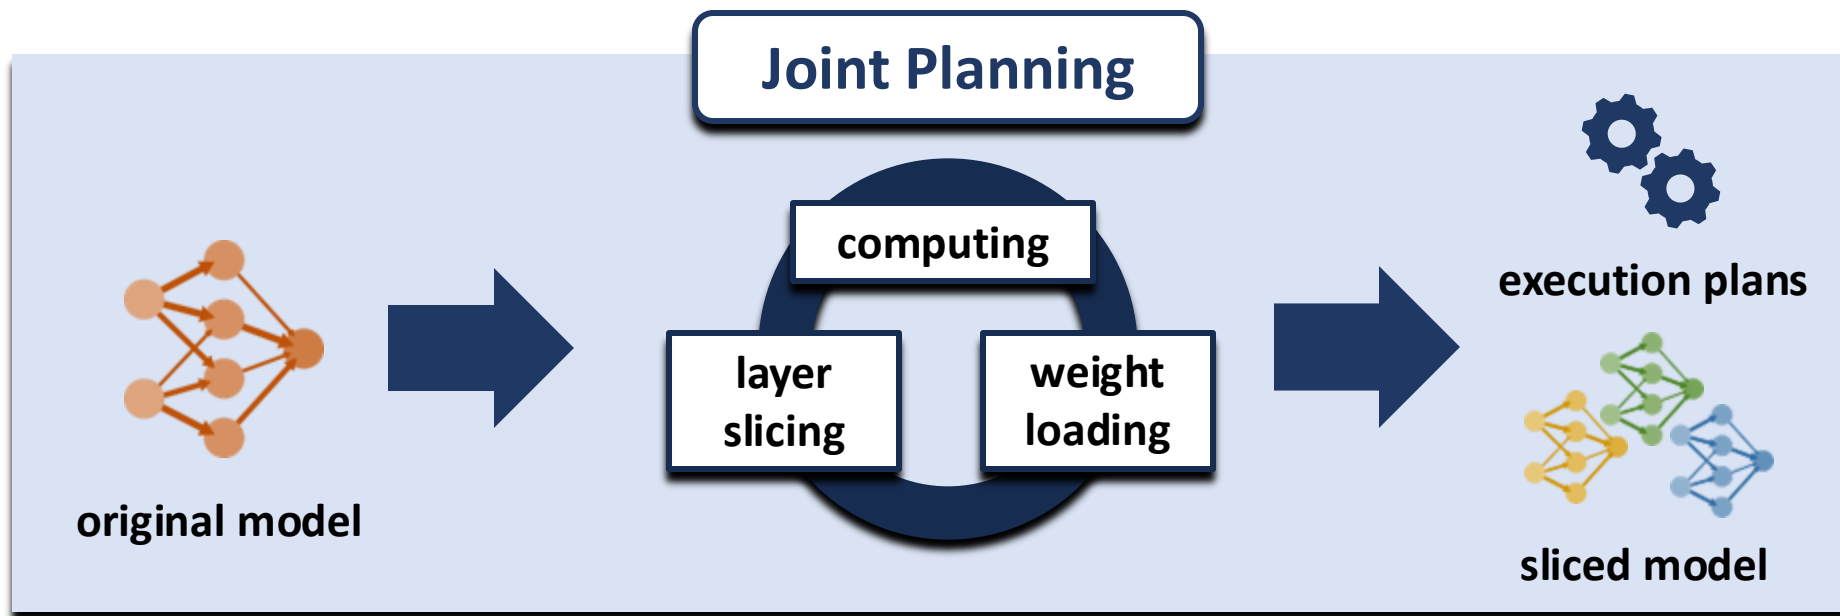
adaptability

Our approaches

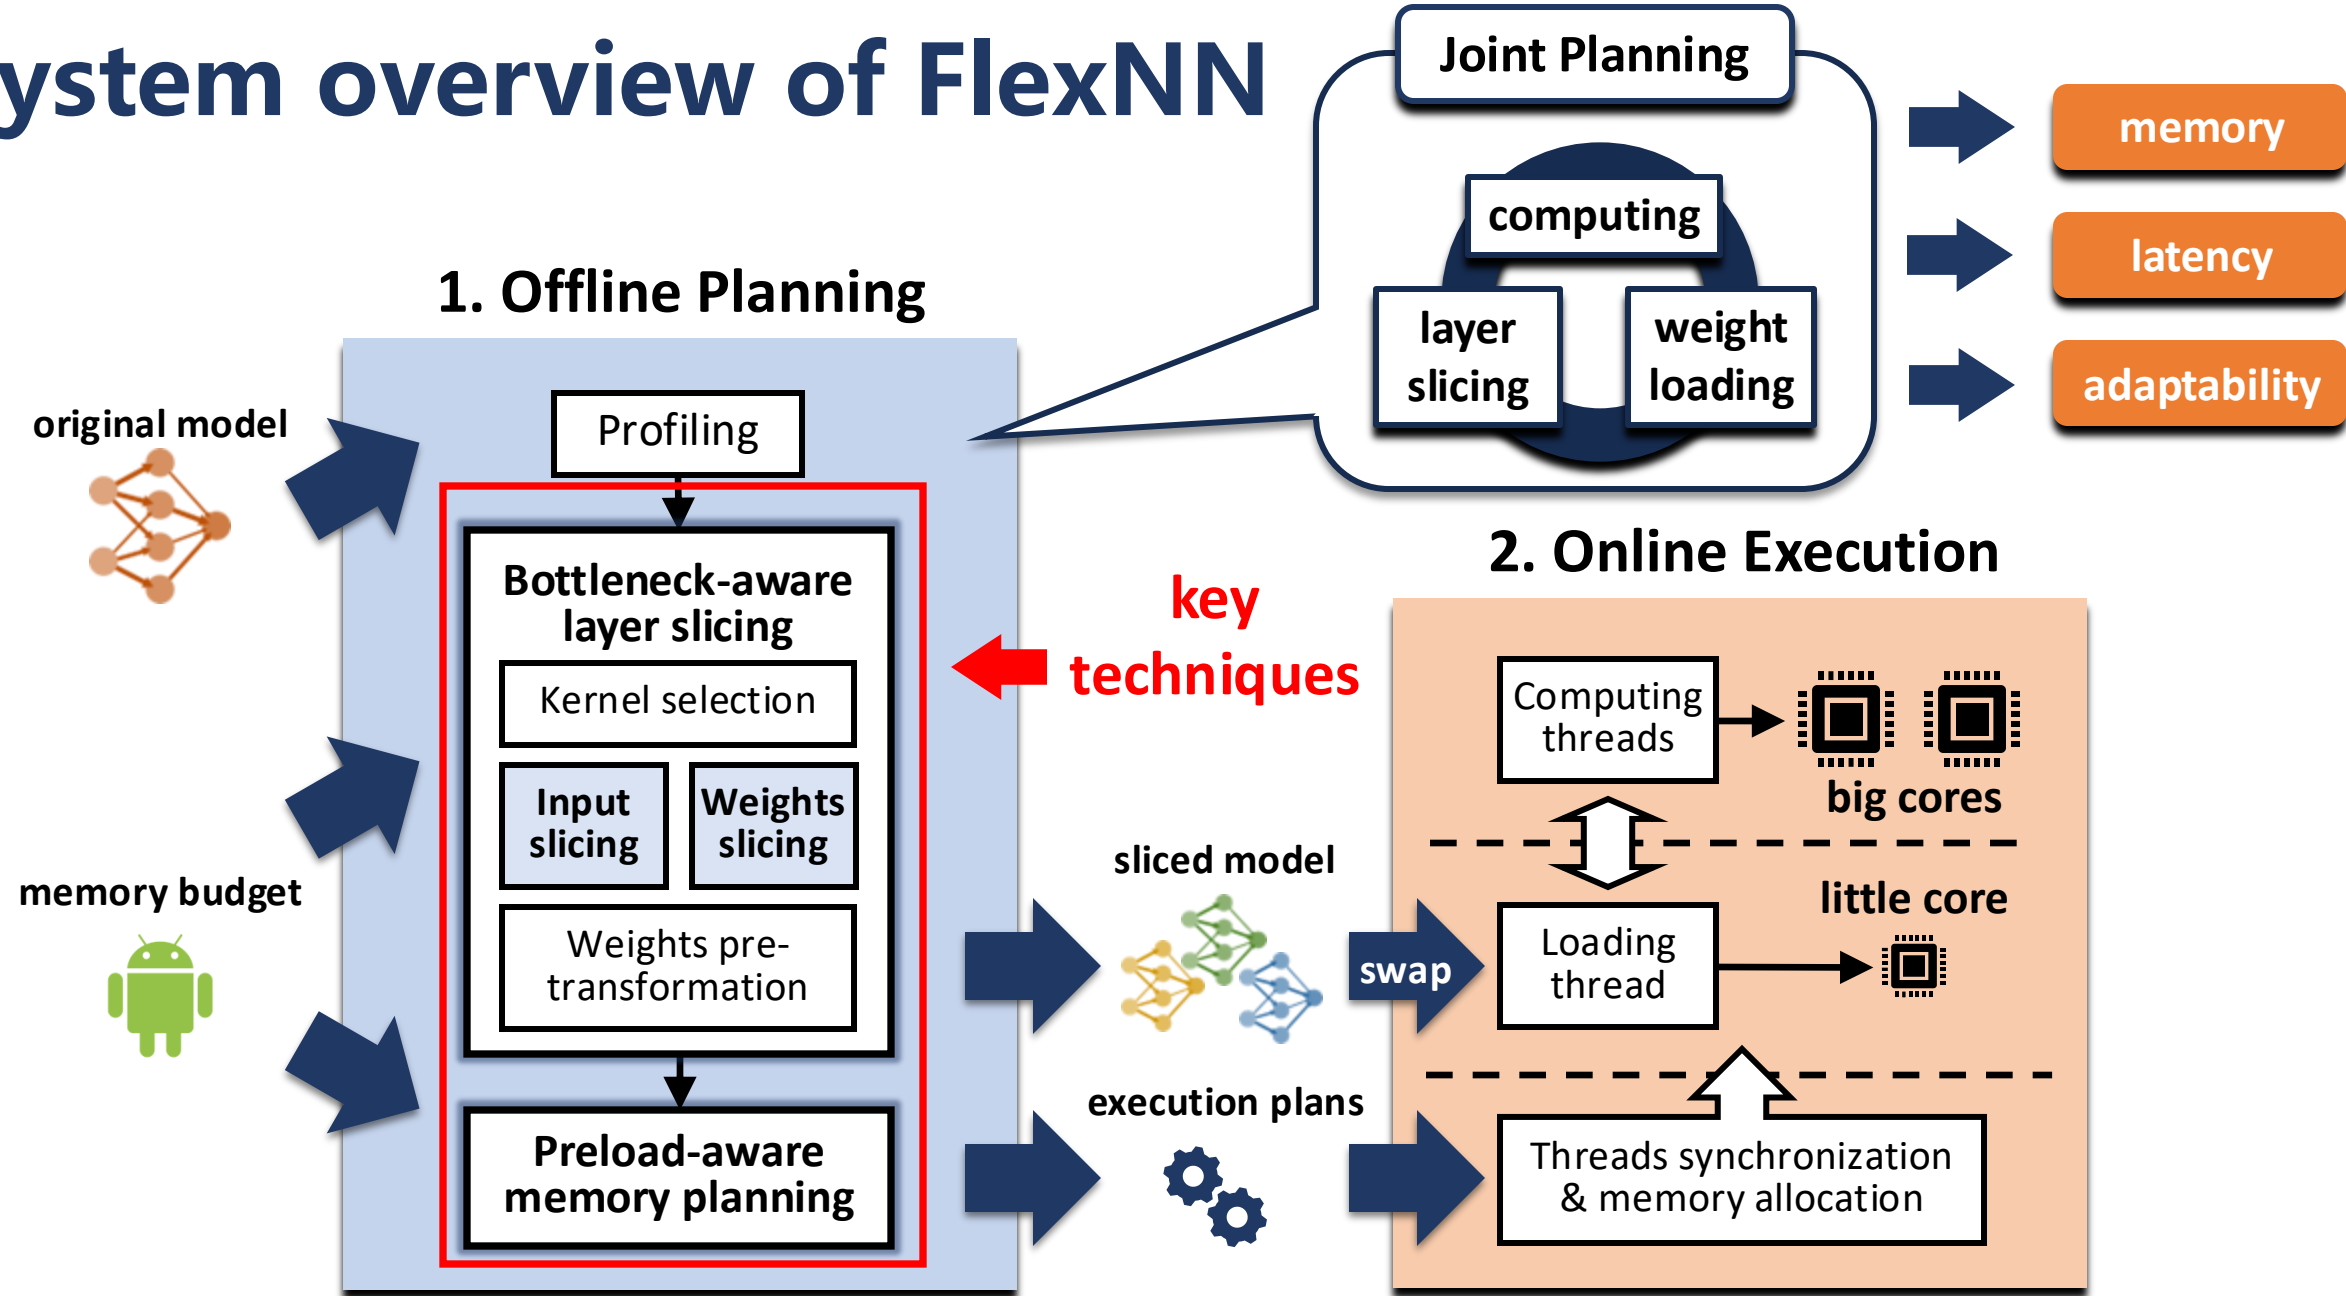
Layer slicing

Offline planning & preparation

Efficient planning algorithm



System overview of FlexNN




Weights slicing: Fully-Connected (FC) example

Implementation: split a large FC layer into sub-layers.

Slicing strategy: *maximize* the slice size below the memory budget.


- Reduce number of layers to schedule and run.



Weight slicing example of 8x6 FC.

The weights are partitioned into 6 slices to save 83% of memory usage.

Fully-connected (FC) layers involve simple linear operations (e.g., matrix multiplications).



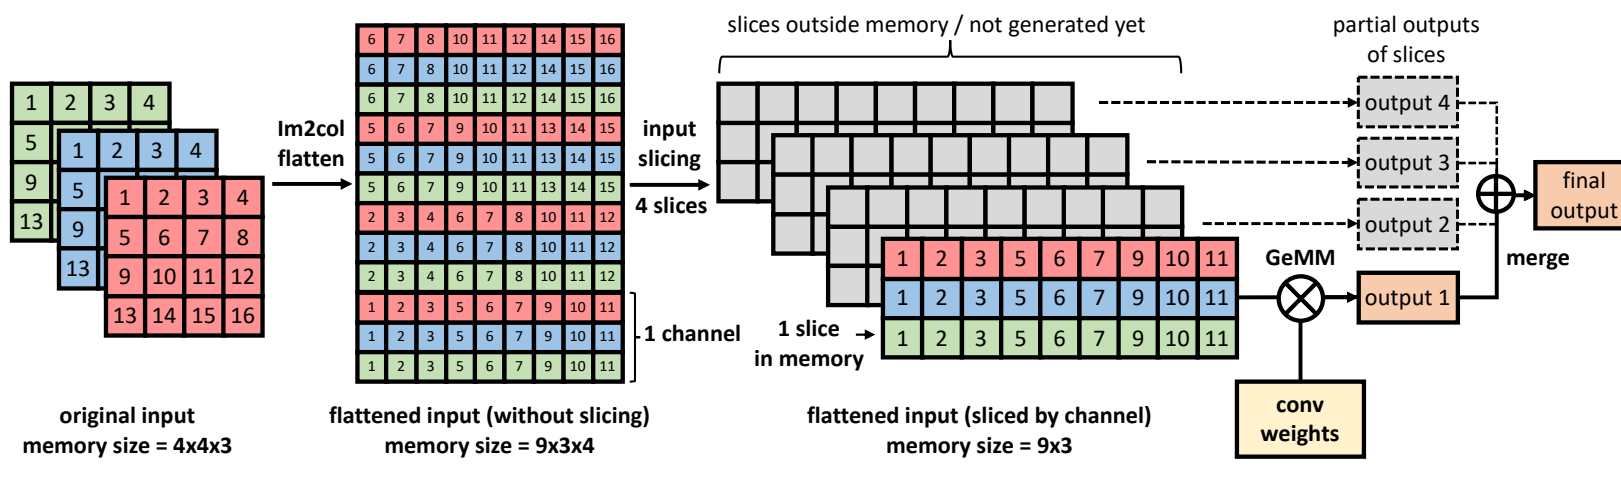
Layer-wise latency-memory trade-off of weight slicing.

Input slicing: Im2col + GeMM example

Implementation: the GeMM input is flattened and calculated slice by slice.

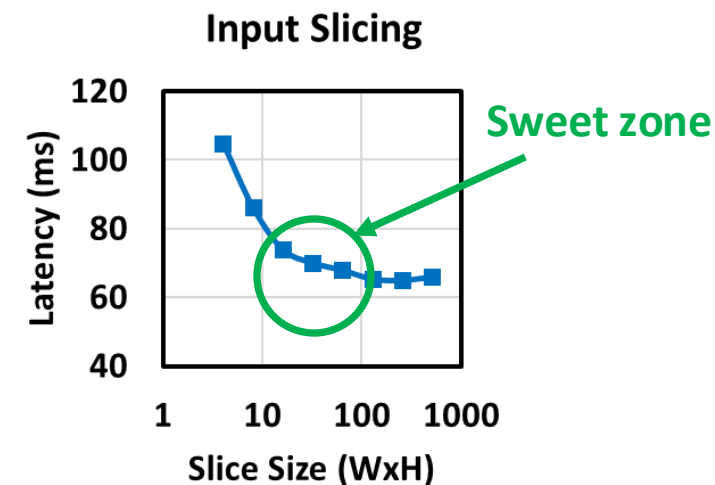
Slicing strategy: *minimize* the slice size above a certain threshold.

- Memory saving with negligible overhead.
- Cannot fully utilize parallel acceleration (e.g., SIMD) if the slice size is too small.



Input slicing example of Im2col+GeMM 3x3 Conv.

The Im2col-flattened input is partitioned into 4 slices to save 75% of memory usage.



Layer-wise latency-memory trade-off of input slicing.

Im2col + GeMM is a commonly used Conv kernel, which involves flattening the input tensor (*i.e.*, Im2col), and multiplication to the weights (*i.e.*, GeMM).

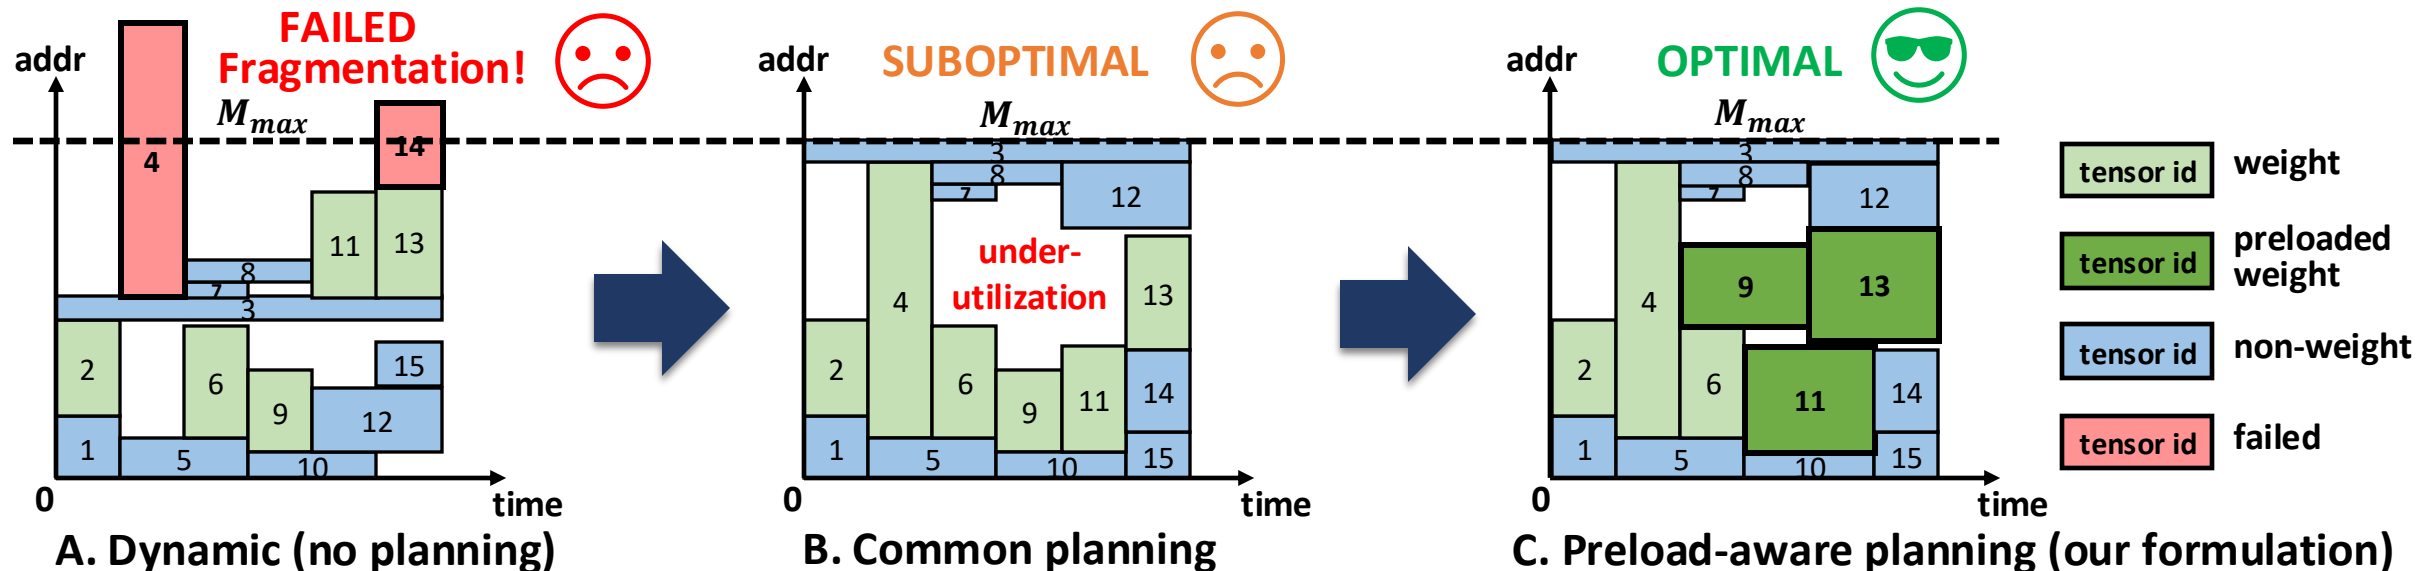
Preload-Aware Memory Planning

Preliminary: 2DBP (2D Bin Packing, NP-Hard) formulation of memory planning.

- Allocating tensors (space & time) is equal to placing rectangles in a plane.
- Mapping: $\text{time}^* \rightarrow x$ value; memory address $\rightarrow y$ value

Our case: a more complex 2DBP variant which is aware of weight preloading.

- Difference: the allocation time of weight tensors is planning output, not input.
- **Concern: significant planning cost when adapting to a new memory budget.**



* Logical time in this case refers to the layer index, since the runtime scheduling is in the granularity of layers.

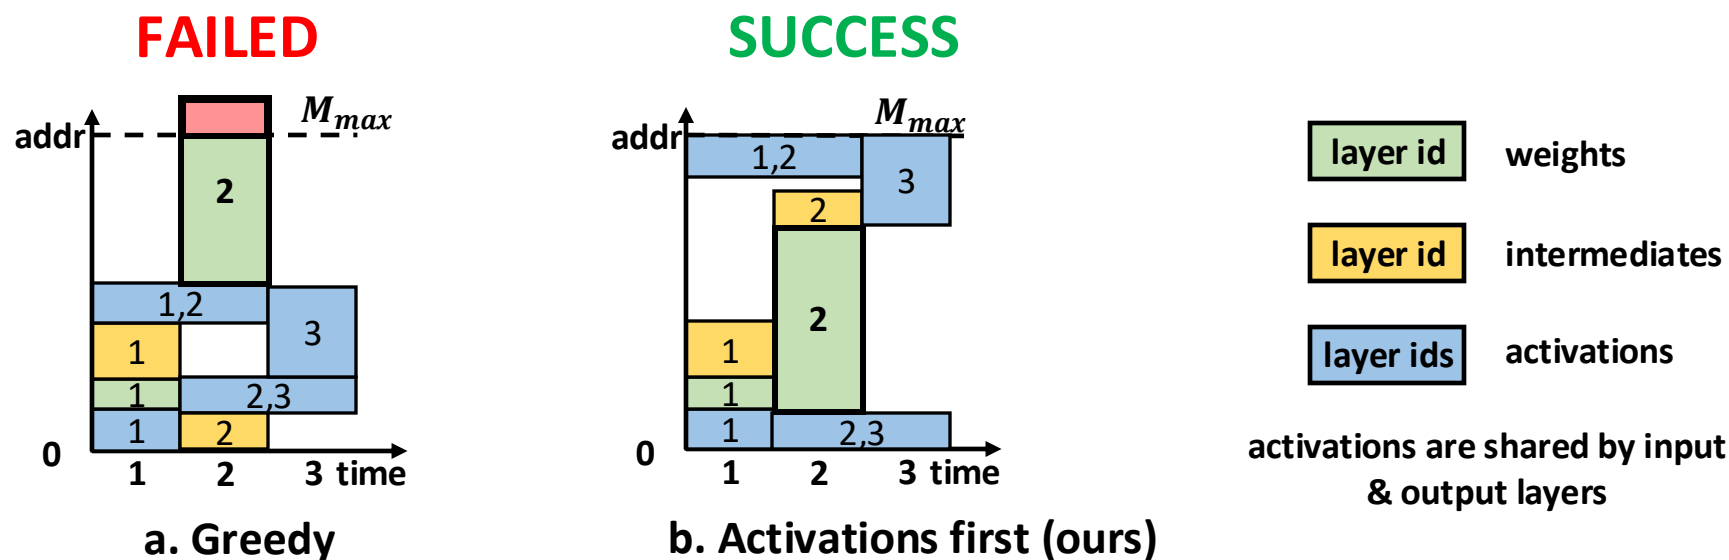
Our approach: prioritize activations

Insight: fragments are caused by *long-lifecycle* tensors.

- *Activations* have longer lifecycle than the others during inference.

Solution: plan *activations* first.

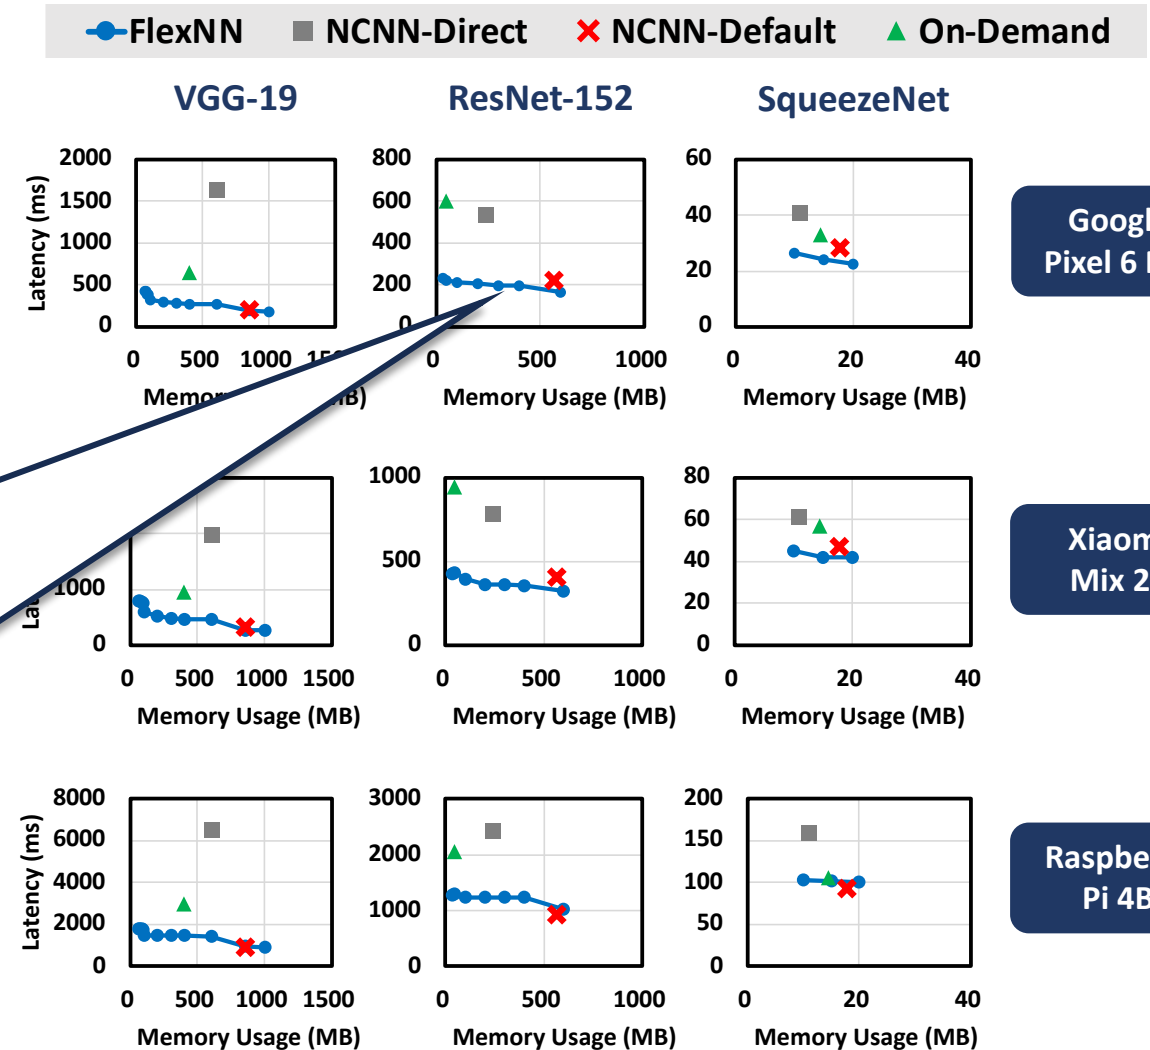
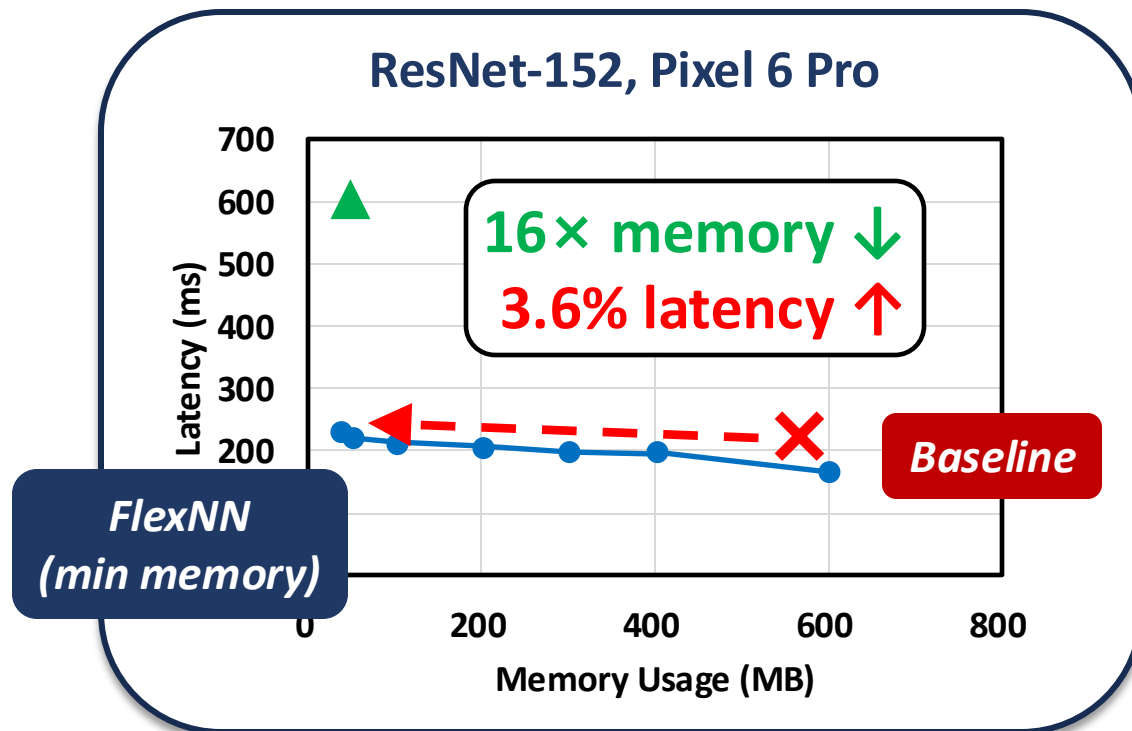
1. Plan all the activations in the model.
2. Layer-wisely plan weights and intermediates.



Comparison between greedy and activation first planning.

Implementation & end-to-end evaluation

- **Implementation:** atop *NCNN*^[1] with 12.3k LoC added (support ARMv8 CPUs).
 - Code: <https://github.com/xxxxyu/FlexNN>.
- **Evaluation:** covers *6 models* and *3 devices* of different types.



End-to-end latency and memory results (partial).
FlexNN achieves better latency-memory trade-offs.

[1] Hui Ni, and The ncnn contributors. Ncnn. 2017, <https://github.com/Tencent/ncnn>.

Adaptability evaluation under changing budgets

FlexNN is able to adapt to new memory budget in ~ 1s.


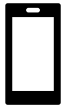






Figure 3. Real-time latency and memory under changing memory budgets.

Summary on FlexNN



-  **Scenario:** memory-constrained DNN inference on edge devices.
-  **Goal:** fast adaption to the given memory budget, with acceptable latency overhead.
-  **Key Idea:** streaming inference with slicing-loading-computing joint planning.
-  **Techniques:** bottleneck-aware layer slicing & preload-aware memory planning.
-  **Key Results:** 16× memory ↓ with 3.6% latency ↑, and ~ 1s adaption.

Thanks!



<https://github.com/xxxxyu/FlexNN>



lixiangy22@mails.tsinghua.edu.cn